

SUMMARY

AI Engineering student with demonstrated depth in Generative AI, RAG architecture, and Agentic systems — LangChain, LangGraph, MCP, and VectorDB. Six live applications and five shipped ML models.

TECHNICAL SKILLS

Agentic AI & LLMs: LangChain, LangGraph, MCP, Multi-LLM Orchestration, Prompt Engineering
Model Development: Scikit-learn, TensorFlow, Keras, Fine-Tuning, PEFT, LoRA
Data & Memory Systems: RAG Pipelines, ChromaDB, FAISS, MySQL, MongoDB
Programming Languages: Python, Java, C++, SQL
Deployment & Automation: Streamlit, Render, Docker, n8n
Scientific Computing & Visualization: Pandas, NumPy, SciPy, Matplotlib, Seaborn, Plotly
Soft Skills: Complex Systems Thinking, Adaptability, Clear Technical Communication, Critical Reasoning

PROJECTS

Natural Language Data Analytics Engine <i>Python LangChain LLaMA 4 Scout Groq API SQLite3 Plotly Streamlit</i> Built an AI-powered data intelligence platform that interprets conversational queries and autonomously orchestrates structured database operations for non-technical users. - Engineered dual-layer query sanitization enforcing strict read-only execution at both prompt and runtime layers. - Implemented multi-format data ingestion supporting CSV, Excel, JSON, and SQLite with Plotly-powered auto-visualization. - Developed a context-aware query suggestion engine that generates intelligent question recommendations based on the uploaded dataset schema. Agentic Chatbot with Real-Time Web Intelligence <i>Python LangChain LangGraph Tavily API FastAPI Streamlit Groq Gemini</i> Architected a production-grade AI agent with live internet access, deployed across a decoupled FastAPI backend and Streamlit frontend on independent cloud platforms. - Integrated Tavily API for real-time web retrieval, grounding agent responses in current information beyond static model knowledge. - Implemented runtime model switching between Groq Llama 3.3 70B and Gemini 2.5 Flash Lite with configurable system prompts and toggleable web search. - Deployed across two independent platforms — Render for backend, Streamlit Cloud for frontend — with a 5-query conversation memory window. AI-Driven Personalized Job Outreach Generator <i>Python LangChain MetaLlama ChromaDB BeautifulSoup Streamlit</i> Engineered a recruitment automation system that ingests a job posting URL and generates a hyper-personalized cold email in seconds. - Implemented automated job description extraction via web scraping and prompt engineering to accurately distill role requirements. - Integrated ChromaDB semantic retrieval to surface contextually relevant portfolio projects for every unique job description. - Designed to produce non-generic, human-quality outreach enabling job seekers to bypass saturated applicant pipelines at scale. Rain Tomorrow Prediction <i>Python Scikit-learn Random Forest Pandas OpenWeatherMap API Streamlit</i> Developed a meteorological prediction engine trained on 10,000+ records achieving 83.5% accuracy with real-time city-level forecasting. - Trained and optimized a Random Forest classifier through rigorous feature selection and hyperparameter tuning on 10,000+ meteorological records. - Integrated OpenWeatherMap API for live atmospheric data inference, delivering confidence-scored predictions for any queried city. - Deployed as a publicly accessible Streamlit application demonstrating full end-to-end ML pipeline from raw data to live inference.	Try it ■ Try it ■ Try it ■ Try it ■
---	---

EDUCATION

Bachelor of Science in Artificial Intelligence Mehran University of Engineering & Technology 2nd Year CGPA: 3.35 / 4.0	2024 – 2028
---	--------------------

CERTIFICATIONS & EXPERIENCE

Google AI Essentials | Coursera | 2025

NASA Space Apps Challenge | Participant, Hyderabad | 2025

- Collaborated with a team to develop an agriculture-based ML project addressing real-world crop and farming challenges.
- Competed alongside participants from across the region in a globally recognized NASA-organized hackathon.